

KAMALESH KUMAR

413-472-6246 kamalvidya.2002@gmail.com [kkk2002](https://www.linkedin.com/in/kkk2002/) [kkmaleshkumar](https://github.com/kkmaleshkumar) [kkmaleshkumar.github.io](https://github.com/kkmaleshkumar.github.io)

Education

University of Massachusetts Amherst

Master of Science in Computer Science (GPA 4.00/4.00)

Sep. 2024 – May 2026

Amherst, MA

Indian Institute of Technology (IIT) Madras

B.Tech in Civil Engineering, Minor in Artificial Intelligence & Machine Learning

Jul. 2020 – May 2024

Chennai, India

Relevant Coursework: Advanced Machine Learning, Robotics, Systems for Data Science, Reinforcement Learning, Multi-Armed Bandits, Computational Cognition, Non-Linear Optimization.

Publications

- **K. Kumar**, J.-A. Delamer, and J. Hughes, “Breaking free from hand-crafted rewards: A genetic programming framework for end-goal-driven reinforcement learning,” (*in review at AAAI 2026*)
- **K. Kumar**, P. P. Kendre, R. D. Manilal, and R. Muthuganapathy, “Sketchcleangan: A generative network to enhance and correct query sketches for improving 3d cad model retrieval systems,” *Computers & Graphics*, vol. 123, 2024
- P. P. Kendre, **K. Kumar**, S. S. K. Jayasree, S. Rajan, A. Jayan, and R. Muthuganapathy, “Sketchcadgan: A generative approach for completing partially drawn query sketches of engineering shapes to enhance retrieval system performance,” *Computers & Graphics*, vol. 115, 2023

Experience

KLA Corporation

May 2025 – Aug. 2025

Machine Learning Intern

Milpitas, CA

- Leveraging continual learning advances to tackle catastrophic forgetting in deep learning based wafer inspection models.
- Performed initial experiments that lead to a ~71% drop in the defect count and a ~52% improvement in precision.
- Collaborating with cross-functional teams to leverage recent research in gradient-replay, adapter & LoRA fine-tuning.

Mitacs Globalink

May 2024 – Aug. 2024

Reinforcement Learning (RL) Research Intern at the Convergence Lab, St. Francis Xavier University

Antigonish, Canada

- Worked on genetic programming (GP) to discover novel reward functions in RL using the DEAP framework.
- Parallelized training of the PPO algorithm across the population of reward functions in the Compute Canada cluster.
- Showed the ability of GP in discovering better performing reward functions in MuJoCo and other Gym environments.

Paris AI Research Institute

May 2023 – Aug. 2023

Reinforcement Learning (RL) Research Intern at the MILES lab, Université Paris-Dauphine - PSL

Paris, France

- Awarded the Charpak Scholarship by the French Embassy in India to research on adversarial & robust RL.
- Investigated state, action, and kernel perturbing adversaries, and established theoretical equivalences between them.
- Proved connections between optimal transport distance, optimal couplings, and adversarial risk in RL

Nuartin Labs

Jun. 2022 – Aug. 2022

Machine Learning Intern

Bangalore, India

- Built pipeline for efficient watermark removal for downstream consumers alongside the founding member.
- Developed an end-to-end pipeline consisting of calibrated localization, image super-resolution, and object segmentation.
- Parallelized the training across four Nvidia T4 GPUs, and deployed final pipeline for clients such as Walmart.

Projects

Continual Reinforcement Learning with Average Reward Criterion | UMass Amherst

Feb. 2025 – present

- Investigating non-stationary environments in reset-free, continual RL settings requiring lifelong agent adaptation.
- Seeking theoretical connections with average-reward POMDPs for modeling partial observability in infinite-horizon tasks.

Autonomous Object Following Robot using ROS and DeepSORT | UMass Amherst

Feb. 2025 – May 2025

- Built a ROS-based object-following robot using YOLO-v3 and DeepSORT for real-time tracking and re-identification.
- Designed a Dockerized ROS Noetic environment on Triton enabling CUDA-accelerated inference, and real-time control.

Real-Time Fake News Detection in Articles Using Apache Flink | UMass Amherst

Sep. 2024 – Dec. 2024

- Developed a real-time streaming pipeline with Apache Flink and ONNX-optimized DistilBERT for fake news detection.
- Optimized system performance for throughput, latency, fault tolerance, and resource efficiency in a scalable deployment.

Improving Sketch Queries for Robust Retrieval of 3D CAD Models | IIT Madras

Aug. 2022 – Dec. 2023

- Designed a two-stage cascaded GAN architecture to facilitate sketch completion of incomplete query sketches.
- Proposed a novel three-branch factorization based on conditional Wasserstein Generative Adversarial Network (GAN) to clean defective sketches and thus improvised a dataset of 58K CAD sketches. Published at Computer & Graphics'24.

Technical Skills

Languages: Python, C++, MATLAB, L^AT_EX, C, SQL

Libraries: PyTorch, TensorFlow, Transformers (Huggingface), Stable-baselines, Gymnasium, OpenCV, Numpy, Pandas

Technologies/Frameworks: ROS, Ray, Linux, Git, Spark, Flink, Hadoop, Docker, Azure